



Source Routing Using P4 match-action tables

(<https://github.com/Tigohsr/sourceRouting-P4>)

Felipe Novais

Thiago Henrique S Rodrigues

Tables in Control Plane



table_set_default sroute_exact drop table_set_default ipv4_lpm drop		
table_add sroute_exact sroute_forward 10.0.0.10 10.0.2.10 => 0x000000000000020302 table_add sroute_exact sroute_forward 10.0.2.10 10.0.0.10 => 0x000000000000010101		
table_add sroute_exact sroute_forward 10.0.1.10 10.0.0.10 => 0x000000000000010303 table_add sroute_exact sroute_forward 10.0.1.10 10.0.2.10 => 0x000000000000020301		
table_add ipv4_lpm ipv4_forward 10.0.1.10/32 => 00:04:00:00:00:01 2		
~		
command_s1.txt	2,1	All
table_set_default sroute_exact drop table_set_default ipv4_lpm drop		
table_add sroute_exact sroute_forward 10.0.0.10 10.0.2.10 => 0x000000000000020302 table_add sroute_exact sroute_forward 10.0.2.10 10.0.0.10 => 0x000000000000010101		
table_add sroute_exact sroute_forward 10.0.1.10 10.0.0.10 => 0x000000000000010303 table_add sroute_exact sroute_forward 10.0.1.10 10.0.2.10 => 0x000000000000020301		
table_add ipv4_lpm ipv4_forward 10.0.1.10/32 => 00:04:00:00:00:01 2		
~		
command_s2.txt	4,1	All
table_set_default sroute_exact drop table_set_default ipv4_lpm drop		
table_add sroute_exact sroute_forward 10.0.0.10 10.0.2.10 => 0x000000000000020302 table_add sroute_exact sroute_forward 10.0.2.10 10.0.0.10 => 0x000000000000010101		
table_add sroute_exact sroute_forward 10.0.1.10 10.0.0.10 => 0x000000000000010303 table_add sroute_exact sroute_forward 10.0.1.10 10.0.2.10 => 0x000000000000020301		
table_add ipv4_lpm ipv4_forward 10.0.1.10/32 => 00:04:00:00:00:01 1		
command_s3.txt	10,1	All

Defines



```
7
8 #define MAX_HOPS 9
9 // H1 - 10.0.0.10
10 #define H1 0x0a00000a
11 // H2 - 10.0.1.10
12 #define H2 0x0a00010a
13 // H3 - 10.0.2.10
14 #define H3 0x0a00020a
15
```

Parsers



```
73
74     state parse_ethernet {
75         packet.extract(hdr.ethernet);
76         transition select(hdr.ethernet.etherType) {
77             TYPE_SRCROUTING: parse_srcRouting;
78             TYPE_IPV4: parse_ipv4;
79             default: accept;
80         }
81     }
82
83     state parse_srcRouting {
84         packet.extract(hdr.srcRoutes.next);
85         transition select(hdr.srcRoutes.last.bos) {
86             1: parse_ipv4;
87             default: parse_srcRouting;
88         }
89     }
90
91     state parse_ipv4 {
92         packet.extract(hdr.ipv4);
93         transition accept;
94     }
95
```

Tables in P4

```
177     table ipv4_lpm {
178         key = {
179             hdr.ipv4.dstAddr: lpm;
180         }
181         actions = {
182             ipv4_forward;
183             drop;
184             NoAction;
185         }
186         size = 1024;
187         default_action = NoAction();
188     }
189
190     table sroute_exact {
191         key = {
192             hdr.ipv4.srcAddr: exact;
193             hdr.ipv4.dstAddr: exact;
194         }
195         actions = {
196             sroute_forward;
197             drop;
198             NoAction;
199         }
200         size = 1024;
201         default_action = NoAction();
202     }
```


Apply

```
198  apply {
199      // If it's H2, uses destination routing | else if it's etherType is srcRouting use source routing
200      if (hdr.ipv4.dstAddr == H2) {
201          ipv4_lpm.apply();
202      }
203      } else if (hdr.ethernet.etherType == TYPE_SRCROUTING) {
204          // Check if it's an empty source routing header, then initialize using table values
205          if (hdr.srcRoutes[0].port == 0 && hdr.srcRoutes[0].bos == 1) {
206              hdr.srcRoutes.push_front(2);
207              sroute_exact.apply();
208          }
209
210          // Usual flow for source routing used in task 3
211          if (hdr.srcRoutes[0].isValid()){
212              if (hdr.srcRoutes[0].bos == 1) {
213                  srcRoute_finish();
214              }
215              srcRoute_nhop();
216              if (hdr.ipv4.isValid()){
217                  update_ttl();
218              }
219          }
220      }
221      else {
222          drop();
223      }
224  }
```

Action ipv4_forward



```
133  action ipv4_forward(macAddr_t dstAddr, egressSpec_t port) {  
134      standard_metadata.egress_spec = port;  
135      hdr.ethernet.srcAddr = hdr.ethernet.dstAddr;  
136      hdr.ethernet.dstAddr = dstAddr;  
137      hdr.ipv4.ttl = hdr.ipv4.ttl - 1;  
138  }
```

Action sroute_forward



```
140  action sroute_forward(bit<72> sourceRoute) {
141      // Initializing bos with zeroes for all elements in srcRoutes
142      hdr.srcRoutes[0].bos = 0;
143      hdr.srcRoutes[1].bos = 0;
144      hdr.srcRoutes[2].bos = 0;
145      hdr.srcRoutes[3].bos = 0;
146      hdr.srcRoutes[4].bos = 0;
147      hdr.srcRoutes[5].bos = 0;
148      hdr.srcRoutes[6].bos = 0;
149      hdr.srcRoutes[7].bos = 0;
150      hdr.srcRoutes[8].bos = 1;
151
152      // Decoding values from table using bitmask technique
153      hdr.srcRoutes[0].port = (bit<15>) ((sourceRoute & 0x0000000000000000ff));
154      hdr.srcRoutes[1].port = (bit<15>) ((sourceRoute & 0x00000000000000ff00) >> (8 * 1));
155      hdr.srcRoutes[2].port = (bit<15>) ((sourceRoute & 0x000000000000ff0000) >> (8 * 2));
156      hdr.srcRoutes[3].port = (bit<15>) ((sourceRoute & 0x0000000000ff000000) >> (8 * 3));
157      hdr.srcRoutes[4].port = (bit<15>) ((sourceRoute & 0x00000000ff00000000) >> (8 * 4));
158      hdr.srcRoutes[5].port = (bit<15>) ((sourceRoute & 0x000000ff0000000000) >> (8 * 5));
159      hdr.srcRoutes[6].port = (bit<15>) ((sourceRoute & 0x0000ff000000000000) >> (8 * 6));
160      hdr.srcRoutes[7].port = (bit<15>) ((sourceRoute & 0x00ff00000000000000) >> (8 * 7));
161      hdr.srcRoutes[8].port = (bit<15>) ((sourceRoute & 0xff0000000000000000) >> (8 * 8));
162
163      // Setting the bos base on the next element port
164      if (hdr.srcRoutes[1].port == 0) hdr.srcRoutes[0].bos = 1;
165      if (hdr.srcRoutes[2].port == 0) hdr.srcRoutes[1].bos = 1;
166      if (hdr.srcRoutes[3].port == 0) hdr.srcRoutes[2].bos = 1;
167      if (hdr.srcRoutes[4].port == 0) hdr.srcRoutes[3].bos = 1;
168      if (hdr.srcRoutes[5].port == 0) hdr.srcRoutes[4].bos = 1;
169      if (hdr.srcRoutes[6].port == 0) hdr.srcRoutes[5].bos = 1;
170      if (hdr.srcRoutes[7].port == 0) hdr.srcRoutes[6].bos = 1;
171      if (hdr.srcRoutes[8].port == 0) hdr.srcRoutes[7].bos = 1;
172
173      // The last element will always have bos 1
174      hdr.srcRoutes[8].bos = 1;
175  }
```


send.py



```
34 def main():
35     if len(sys.argv)<2:
36         print 'pass 2 arguments: <destination>'
37         exit(1)
38
39     addr = socket.gethostbyname(sys.argv[1])
40     iface = get_if()
41     print "sending on interface %s to %s" % (iface, str(addr))
42
43     pkt = Ether(src=get_if_hwaddr(iface), dst='ff:ff:ff:ff:ff:ff');
44     pkt.show2()
45     if ("10.0.1.10" != sys.argv[1]):
46         pkt = pkt / SourceRoute(bos=1, port=int(0))
47
48     pkt = pkt / IP(src=get_if_addr(iface), dst=addr) / UDP(dport=4321, sport=1234)
49     pkt.show2()
50     sendp(pkt, iface=iface, verbose=False)
51
52
53
54
55 if __name__ == '__main__':
56     main()
```

Bitmask explanation



Consider we have 0x0201:

0x0201 = 0000 0010 0000 0001 <- Binary representation

0x0201 & 0x00FF = 0000 0010 0000 0001 & 0000 0000 1111 1111

bitwise and

```
0000 0010 0000 0001
0000 0000 1111 1111
-----
0000 0010 0000 0001 = 0x0001
```

0x0201 & 0xFF00 = 0000 0010 0000 0001 & 1111 1111 0000 0000

bitwise and

```
0000 0010 0000 0001
1111 1111 0000 0000
-----
0000 0010 0000 0000 = 0x0200
```

Doing a right shift by 8 bits we isolate the second 2 bytes.

0x0200 >> 8 = 0000 0010 0000 0000 = 0000 00010 = 0x02

right shift



Thank you!